

DeepLearning.AI TensorFlow Developer

Contents

1	Introduction to TensorFlow for AI, ML, and DL	1
1.0.1	Callbacks	1
1.0.2	Upload Custom Images	1
1.0.3	ImageDataGenerator	2
2	CNN's in TensorFlow	3
2.0.1	Using OS and Splitting Data	3
2.0.2	Building a CNN with IDG	4
2.0.3	Transfer Learning: Built-in Models	4
2.0.4	Reading Images from CSVs	5
2.0.5	Multi-Class Classification	6
3	Natural Language Processing	7
3.0.1	7

1 Introduction to TensorFlow for AI, ML, and DL

1.0.1 Callbacks

We can use **callbacks** in order to stop training when we reach a certain accuracy we desire. This is to stop the loss from beginning to increase again if we start to overfit the model. [Click here](#) to see the TensorFlow Callbacks documentation.

```
1 import tensorflow as tf
2 print(tf.__version__)
3
4 class myCallback(tf.keras.callbacks.Callback):
5     def on_epoch_end(self, epoch, logs={}):
6         if(logs.get('accuracy')>0.6): # might need to use 'acc' instead
7             print("\nReached 60% accuracy so cancelling training!")
8             self.model.stop_training = True
9
10 callbacks = myCallback()
11
12 mnist = tf.keras.datasets.fashion_mnist
13 (x_train, y_train), (x_test, y_test) = mnist.load_data()
14 x_train, x_test = x_train / 255.0, x_test / 255.0
15
16 model = tf.keras.models.Sequential([
17     tf.keras.layers.Flatten(),
18     tf.keras.layers.Dense(512, activation=tf.nn.relu),
19     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
20 ])
21
22 model.compile(optimizer='adam',
23               loss='sparse_categorical_crossentropy',
24               metrics=['accuracy'])
25 model.fit(x_images, y_labels, epochs=10, callbacks=[callbacks])
```

1.0.2 Upload Custom Images

We can use the below code to upload a custom image and use it on a trained model.

```
1 import numpy as np
2 from google.colab import files
3 from keras.preprocessing import image
4
5 uploaded = files.upload()
6
7 for fn in uploaded.keys():
8     # predicting images
9     path = '/content/' + fn
10    img = image.load_img(path, target_size=(300, 300))
11    x = image.img_to_array(img)
12    x = np.expand_dims(x, axis=0)
13
14    images = np.vstack([x])
15    classes = model.predict(images, batch_size=10)
16    print(classes[0])
17    if classes[0]>0.5:
18        print(fn + " is a human")
19    else:
20        print(fn + " is a horse")
```

1.0.3 ImageDataGenerator

```
1 import tensorflow as tf
2 import os
3 import zipfile
4 from os import path, getcwd, chdir
5 from tensorflow.keras.optimizers import RMSprop
6     from tensorflow.keras.preprocessing.image import ImageDataGenerator
7
8 # Import and extract zip file containing images
9 path = f"{getcwd()}/../tmp2/happy-or-sad.zip"
10 zip_ref = zipfile.ZipFile(path, 'r')
11 zip_ref.extractall("/tmp/h-or-s")
12 zip_ref.close()
13
14 def train_happy_sad_model():
15     DESIRED_ACCURACY = 0.999
16
17     class myCallback(tf.keras.callbacks.Callback):
18         def on_epoch_end(self, epoch, logs={}):
19             if(logs.get('acc')>DESIRED_ACCURACY):
20                 print('\nReached 100% accuracy so stopping training.')
21                 self.model.stop_training = True
22
23     callbacks = myCallback()
24
25     # Define and Compile the Model.
26     model = tf.keras.models.Sequential([
27         tf.keras.layers.Conv2D(64, (3,3), activation='relu', input_shape=(150,150,3)),
28         tf.keras.layers.MaxPooling2D(2,2),
29         tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
30         tf.keras.layers.MaxPooling2D(2,2),
31         tf.keras.layers.Conv2D(16, (3,3), activation='relu'),
32         tf.keras.layers.MaxPooling2D(2,2),
33         tf.keras.layers.Flatten(),
34         tf.keras.layers.Dense(512),
35         tf.keras.layers.Dense(1, activation='sigmoid')
36     ])
37
38     model.compile(optimizer=RMSprop(lr=0.001),
39                 loss='binary_crossentropy',
40                 metrics=['accuracy'])
41
42     # Create an instance of an ImageDataGenerator
43     train_datagen = ImageDataGenerator(rescale=1./255)
44
45     train_generator = train_datagen.flow_from_directory(
46         '/tmp/h-or-s', # directory containing the images
47         target_size=(150,150),
48         class_mode='binary'
49     )
50
51     history = model.fit(
52         train_generator, # data generator object
53         epochs=50,
54         callbacks=[callbacks],
55         verbose=1
56     )
57
58     return history.history['acc'][-1]
```

2 CNN's in TensorFlow

2.0.1 Using OS and Splitting Data

```
1 path_cats_and_dogs = f"{getcwd()}/../tmp2/cats-and-dogs.zip"
2 shutil.rmtree('/tmp')
3
4 local_zip = path_cats_and_dogs
5 zip_ref = zipfile.ZipFile(local_zip, 'r')
6 zip_ref.extractall('/tmp')
7 zip_ref.close()
8
9 try: # Make directories for training and testing, along with class subdirectories
10     os.mkdir('/tmp/cats-v-dogs/')
11     os.mkdir('/tmp/cats-v-dogs/training/')
12     os.mkdir('/tmp/cats-v-dogs/testing/')
13     os.mkdir('/tmp/cats-v-dogs/training/cats/')
14     os.mkdir('/tmp/cats-v-dogs/testing/cats/')
15     os.mkdir('/tmp/cats-v-dogs/training/dogs/')
16     os.mkdir('/tmp/cats-v-dogs/testing/dogs/')
17 except OSError:
18     pass
19
20 def split_data(SOURCE, TRAINING, TESTING, SPLIT_SIZE):
21     """
22     a SOURCE directory containing the files
23     a TRAINING directory that a portion of the files will be copied to
24     a TESTING directory that a portion of the files will be copied to
25     a SPLIT SIZE to determine the portion
26     """
27     source_list = os.listdir(SOURCE) # get list of files
28     random.sample(source_list, len(source_list)) # shuffle the list
29     train_images = source_list[:int(len(source_list)*SPLIT_SIZE)]
30     testing_images = source_list[int(len(source_list)*SPLIT_SIZE):]
31
32     for img in train_images:
33         if os.path.getsize(SOURCE+img) != 0: # make sure not empty file
34             copyfile(SOURCE+img, TRAINING+img)
35
36     for img in testing_images:
37         if os.path.getsize(SOURCE+img) != 0: # make sure not empty file
38             copyfile(SOURCE+img, TESTING+img)
39
40     return None
41
42 CAT_SOURCE_DIR = "/tmp/PetImages/Cat/"
43 TRAINING_CATS_DIR = "/tmp/cats-v-dogs/training/cats/"
44 TESTING_CATS_DIR = "/tmp/cats-v-dogs/testing/cats/"
45 DOG_SOURCE_DIR = "/tmp/PetImages/Dog/"
46 TRAINING_DOGS_DIR = "/tmp/cats-v-dogs/training/dogs/"
47 TESTING_DOGS_DIR = "/tmp/cats-v-dogs/testing/dogs/"
48
49 split_size = .9
50 split_data(CAT_SOURCE_DIR, TRAINING_CATS_DIR, TESTING_CATS_DIR, split_size)
51 split_data(DOG_SOURCE_DIR, TRAINING_DOGS_DIR, TESTING_DOGS_DIR, split_size)
```

2.0.2 Building a CNN with IDG

```
1 model = tf.keras.models.Sequential([
2     tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(150,150,3)),
3     tf.keras.layers.MaxPooling2D(2,2),
4     tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
5     tf.keras.layers.MaxPooling2D(2,2),
6     tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
7     tf.keras.layers.MaxPooling2D(2,2),
8     tf.keras.layers.Flatten(),
9     tf.keras.layers.Dense(512, activation='relu'),
10    tf.keras.layers.Dense(1, activation='sigmoid')
11 ])
12
13 model.compile(optimizer=RMSprop(lr=0.001), loss='binary_crossentropy',
14               metrics=['acc'])
15
16 # Feed training data into our IDG object
17 TRAINING_DIR = '/tmp/cats-v-dogs/training/'
18 train_datagen = ImageDataGenerator(rescale=1./255,
19                                     rotation_range=45,
20                                     horizontal_flip=True,
21                                     vertical_flip=True,
22                                     shear_range=0.2,
23                                     zoom_range=0.2,
24                                     fill_mode='nearest')
25
26 train_generator = train_datagen.flow_from_directory(TRAINING_DIR,
27                                                     target_size=(150,150),
28                                                     batch_size=10,
29                                                     class_mode='binary')
30
31 # Feed testing data into our IDG object
32 VALIDATION_DIR = '/tmp/cats-v-dogs/testing/'
33 validation_datagen = ImageDataGenerator(rescale=1./255)
34
35 validation_generator = validation_datagen.flow_from_directory(VALIDATION_DIR,
36                                                             target_size=(150,150),
37                                                             batch_size=10,
38                                                             class_mode='binary')
39
40 history = model.fit_generator(train_generator,
41                               epochs=2,
42                               verbose=1,
43                               validation_data=validation_generator)
```

2.0.3 Transfer Learning: Built-in Models

```
1 path_inception = f"{getcwd()}/../tmp2/"
2     inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5"
3
4 # Import the inception model
5 from tensorflow.keras.applications.inception_v3 import InceptionV3
6
7 # Create an instance of the inception model from the local pre-trained weights
8 local_weights_file = path_inception
9
10 pre_trained_model = InceptionV3(input_shape=(150,150,3), include_top=False,
11                                 weights=None)
12
13 pre_trained_model.load_weights(local_weights_file)
```

```

1 # Make all the layers in the pre-trained model non-trainable
2 for layer in pre_trained_model.layers:
3     layer.trainable = False
4
5 last_layer = pre_trained_model.get_layer('mixed7')
6 print('last layer output shape: ', last_layer.output_shape) # (None, 7, 7, 768)
7 last_output = last_layer.output
8
9 # Flatten the output layer to 1 dimension (with input from last pretrained layers)
10 x = layers.Flatten()(last_output)
11
12 x = layers.Dense(1024, activation='relu')(x)
13 x = layers.Dropout(0.2)(x)
14 x = layers.Dense(1, activation='sigmoid')(x)
15
16 model = Model(pre_trained_model.input, x)
17
18 model.compile(optimizer = RMSprop(lr=0.0001),
19               loss = 'binary_crossentropy',
20               metrics = ['acc'])

```

2.0.4 Reading Images from CSVs

```

1 def get_data(filename):
2     """
3     Read the file passed into the function. The first line contains
4     the header (so skip it). Each line contains 785 values, with
5     the first being the label and remaining being the pixel values.
6     You will need to reshape the images into 28x28.
7     """
8     with open(filename) as training_file:
9         labels = []
10        images = []
11        reader = csv.reader(training_file, delimiter = ',')
12        next(reader, None) # skip first line
13
14        for row in reader:
15            labels.append(row[0])
16            images.append(np.array(row[1:]).reshape(28,28))
17
18        labels = np.array(labels).astype(float)
19        images = np.array(images).astype(float)
20        return images, labels
21
22 path_sign_mnist_train = f"{getcwd()}/../tmp2/sign_mnist_train.csv"
23 path_sign_mnist_test = f"{getcwd()}/../tmp2/sign_mnist_test.csv"
24 training_images, training_labels = get_data(path_sign_mnist_train)
25 testing_images, testing_labels = get_data(path_sign_mnist_test)
26
27 print(training_images.shape) # (27455, 28, 28)
28 print(training_labels.shape) # (27455)
29 print(testing_images.shape) # (7172, 28, 28)
30 print(testing_labels.shape) # (7172)
31
32 # Expand dimensions (add 1 to the end)
33 training_images = np.expand_dims(training_images, axis=3) # (27455, 28, 28, 1)
34 testing_images = np.expand_dims(testing_images, axis=3) # (7172, 28, 28, 1)

```

2.0.5 Multi-Class Classification

```
1  # Create an ImageDataGenerator objects
2  train_datagen = ImageDataGenerator(
3      rescale=1./255,
4      rotation_range=45,
5      horizontal_flip=True,
6      vertical_flip=True,
7      zoom_range=0.2 )
8
9  validation_datagen = ImageDataGenerator(rescale=1./255)
10
11 # Create generators (images already loaded, not from directory)
12 train_generator = train_datagen.flow(training_images,
13                                     training_labels,
14                                     batch_size=32)
15
16 valid_generator = validation_datagen.flow(testing_images,
17                                           testing_labels,
18                                           batch_size=32)
19
20 # Define the model
21 model = tf.keras.models.Sequential([
22     tf.keras.layers.Conv2D(64, (3,3), activation='relu', input_shape=(28,28,1)),
23     tf.keras.layers.MaxPooling2D(2,2),
24     tf.keras.layers.Conv2D(128, (3,3), activation='relu'),
25     tf.keras.layers.MaxPooling2D(2,2),
26     tf.keras.layers.Flatten(),
27     tf.keras.layers.Dense(512, activation='relu'),
28     tf.keras.layers.Dense(26, activation='softmax') # 26 classes
29 ])
30
31 # Compile Model.
32 model.compile(optimizer='adam',
33              loss='sparse_categorical_crossentropy',
34              metrics=['acc'])
35
36 # Train the Model
37 history = model.fit_generator(train_generator,
38                              validation_data=valid_generator,
39                              epochs=2,
40                              verbose=1)
41
42 # Access the metrics for plotting
43 acc = history.history['acc']
44 val_acc = history.history['val_acc']
45 loss = history.history['loss']
46 val_loss = history.history['val_loss']
```

3 Natural Language Processing

3.0.1 ...